

# Wasserstein GAN

Input: Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

Task: Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

Method: Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \mathcal{W}(\mathcal{G}_{\theta}, \mu)$$

- $\mathcal{W}$  is the Wasserstein 1-distance, quantifies similarity of probability measures.
- Generator:  $\mathcal{G}_{\theta}$  is a probability distribution on  $X$ .
- Desired property:  $\mathcal{G}_{\hat{\theta}} \approx \mu$ .

Issue:  $\mu$  unknown  $\implies$  not possible to evaluate  $\mathcal{W}(\mathcal{G}_{\theta}, \mu)$ .

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{D \in \text{Lip}(X)} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [D(\mathbf{x})] - \mathbb{E}_{\mathbf{v} \sim \mathcal{G}_{\theta}} [D(\mathbf{v})] \right\} \right\}$$

- **Discriminator:**  $D: X \rightarrow \mathbb{R}$  is a 1-Lipschitz function.
- **Generator:**  $\mathcal{G}_{\theta}$  is a probability distribution on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}} \approx \mu$ .

**Issue:**  $\mu$  unknown  $\implies$  not possible to evaluate  $\mathcal{W}(\mathcal{G}_{\theta}, \mu)$ .

$\implies$  Re-write using the Kantorovich-Rubinstein dual characterization of  $\mathcal{W}$ .

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{D \in \text{Lip}(X)} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [D(\mathbf{x})] - \mathbb{E}_{\mathbf{v} \sim \mathcal{G}_{\theta}} [D(\mathbf{v})] \right\} \right\}$$

- **Discriminator:**  $D: X \rightarrow \mathbb{R}$  is a **1-Lipschitz** function.
- **Generator:**  $\mathcal{G}_{\theta}$  is a probability distribution on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}} \approx \mu$ .

**Issue:** How to ensure  $D$  is 1-Lipschitz?

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{D: X \rightarrow \mathbb{R}} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [D(\mathbf{x})] - \mathbb{E}_{\mathbf{v} \sim \mathcal{G}_{\theta}} [D(\mathbf{v})] + \lambda \mathbb{E}_{\tilde{\mathbf{x}} \sim \tilde{\mu}} \left[ \left( \|\partial D(\tilde{\mathbf{x}})\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $D: X \rightarrow \mathbb{R}$  is a measurable function.
- **Generator:**  $\mathcal{G}_{\theta}$  is a probability distribution on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}} \approx \mu$ .
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) \mathbf{v}$  with  $\epsilon \sim U(0, 1)$ .

**Issue:** How to ensure  $D$  is 1-Lipschitz?

$\implies$  softly enforce this condition by adding a penalty term.

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{D: X \rightarrow \mathbb{R}} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [D(\mathbf{x})] - \mathbb{E}_{\mathbf{v} \sim \mathcal{G}_{\theta}} [D(\mathbf{v})] + \lambda \mathbb{E}_{\tilde{\mathbf{x}} \sim \tilde{\mu}} \left[ \left( \|\partial D(\tilde{\mathbf{x}})\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $D: X \rightarrow \mathbb{R}$  is a **measurable function**.
- **Generator:**  $\mathcal{G}_{\theta}$  is a probability distribution on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}} \approx \mu$ .
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) \mathbf{v}$  with  $\epsilon \sim U(0, 1)$ .

**Issue:** Unfeasible to maximise over all measurable mappings

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [\mathbf{D}_{\phi}(\mathbf{x})] - \mathbb{E}_{\mathbf{v} \sim \mathcal{G}_{\theta}} [\mathbf{D}_{\phi}(\mathbf{v})] + \lambda \mathbb{E}_{\tilde{\mathbf{x}} \sim \tilde{\mu}} \left[ \left( \|\partial \mathbf{D}_{\phi}(\tilde{\mathbf{x}})\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $\mathbf{D}_{\phi}: X \rightarrow \mathbb{R}$  a function parametrised by **neural network**.
- **Generator:**  $\mathcal{G}_{\theta}$  is a probability distribution on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}} \approx \mu$ .
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) \mathbf{v}$  with  $\epsilon \sim U(0, 1)$ .

**Issue:** Unfeasible to maximise over all measurable mappings

$\implies$  Parametrise discriminator by a neural network.

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [D_{\phi}(\mathbf{x})] - \mathbb{E}_{\mathbf{v} \sim \mathcal{G}_{\theta}} [D_{\phi}(\mathbf{v})] + \lambda \mathbb{E}_{\tilde{\mathbf{x}} \sim \tilde{\mu}} \left[ \left( \|\partial D_{\phi}(\tilde{\mathbf{x}})\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \rightarrow \mathbb{R}$  a function parametrised by neural network.
- **Generator:**  $\mathcal{G}_{\theta}$  is a **probability distribution** on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}} \approx \mu$ .
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) \mathbf{v}$  with  $\epsilon \sim U(0, 1)$ .

**Issue:** Difficult to parametrise family of probability distributions on  $X$

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [D_{\phi}(\mathbf{x})] - \mathbb{E}_{\mathbf{z} \sim \sigma} [D_{\phi}(G_{\theta}(\mathbf{z}))] + \lambda \mathbb{E}_{\tilde{\mathbf{x}} \sim \tilde{\mu}} \left[ \left( \|\partial D_{\phi}(\tilde{\mathbf{x}})\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \rightarrow \mathbb{R}$  a function parametrised by neural network.
- **Generator:**  $G_{\theta}: Z \rightarrow X$  a **function** parametrised by **neural network**.
- **Desired property:**  $G_{\hat{\theta}}(\mathbf{z})$  is approximately  $\mu$ -distributed when  $\mathbf{z} \sim \sigma$  ( $\sigma$  known).
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) G_{\theta}(\mathbf{z})$  with  $\epsilon \sim U(0, 1)$ .

**Issue:** Difficult to parametrise family of probability distributions on  $X$

$\implies$  Write generator as deterministic function with random input  $\mathbf{z} \sim \sigma$  with  $\sigma$  known.



# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \mathbb{E}_{\mathbf{x} \sim \mu} [D_{\phi}(\mathbf{x})] - \mathbb{E}_{\mathbf{z} \sim \sigma} [D_{\phi}(G_{\theta}(\mathbf{z}))] + \lambda \mathbb{E}_{\tilde{\mathbf{x}} \sim \tilde{\mu}} \left[ \left( \|\partial D_{\phi}(\tilde{\mathbf{x}})\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \rightarrow \mathbb{R}$  a function parametrised by neural network.
- **Generator:**  $G_{\theta}: Z \rightarrow X$  a function parametrised by neural network.
- **Desired property:**  $G_{\hat{\theta}}(\mathbf{z})$  is approximately  $\mu$ -distributed when  $\mathbf{z} \sim \sigma$  ( $\sigma$  known).
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) G_{\theta}(\mathbf{z})$  with  $\epsilon \sim U(0, 1)$ .

**Issue:**  $\mu$  unknown  $\implies$  not possible to compute  $\mu$ - and  $\tilde{\mu}$ -expectations

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \frac{1}{m} \sum_{i=1}^m [D_{\phi}(x_i)] - \mathbb{E}_{\mathbf{z} \sim \sigma} [D_{\phi}(G_{\theta}(\mathbf{z}))] + \lambda \frac{1}{m} \sum_{i=1}^m \left[ \left( \|\partial D_{\phi}(\tilde{x}_i)\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \rightarrow \mathbb{R}$  a function parametrised by neural network.
- **Generator:**  $G_{\theta}: Z \rightarrow X$  a function parametrised by neural network.
- **Desired property:**  $G_{\hat{\theta}}(\mathbf{z})$  is approximately  $\mu$ -distributed when  $\mathbf{z} \sim \sigma$  ( $\sigma$  known).
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) G_{\theta}(\mathbf{z})$  with  $\epsilon \sim U(0, 1)$ .

**Issue:**  $\mu$  unknown  $\implies$  not possible to compute  $\mu$ - and  $\tilde{\mu}$ -expectations

$\implies$  Use empirical distributions given by  $x_i$  and  $\tilde{x}_i = \epsilon_i x_i + (1 - \epsilon_i) G_{\theta}(z_i)$ .

# Wasserstein GAN

**Input:** Unsupervised data  $x_i \in X$  generated by  $\mathbf{x} \sim \mu$ .

**Task:** Sample from  $\mathbf{x} \sim \mu$  where  $\mu$  is unknown.

**Method:** Learn how to sample from  $\mu$  by solving (Arjovsky et al., 2017)

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \frac{1}{m} \sum_{i=1}^m [D_{\phi}(x_i)] - \mathbb{E}_{\mathbf{z} \sim \sigma} [D_{\phi}(G_{\theta}(\mathbf{z}))] + \lambda \frac{1}{m} \sum_{i=1}^m \left[ \left( \|\partial D_{\phi}(\tilde{x}_i)\| - 1 \right)^2 \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \rightarrow \mathbb{R}$  a function parametrised by neural network.
- **Generator:**  $G_{\theta}: Z \rightarrow X$  a function parametrised by neural network.
- **Desired property:**  $G_{\hat{\theta}}(\mathbf{z})$  is approximately  $\mu$ -distributed when  $\mathbf{z} \sim \sigma$  ( $\sigma$  known).
- $\tilde{\mu}$  probability measure generated by  $\epsilon \mathbf{x} + (1 - \epsilon) G_{\theta}(\mathbf{z})$  with  $\epsilon \sim U(0, 1)$ .

**Conclusion:** Above expression is suitable for deep learning.

# Conditional Wasserstein GAN

**Input:** Supervised data  $(x_i, y_i) \in X \times Y$  generated by  $(\mathbf{x}, \mathbf{y}) \sim \mu$ .

**Task:** Sample from  $(\mathbf{x} \mid \mathbf{y} = y) \sim \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$  where posterior is unknown.

**Method:** Learn how to sample from posterior by solving

$$\hat{\theta} \in \arg \min_{\theta} \mathbb{E}_{\mathbf{y}} \left[ \mathcal{W}(\mathcal{G}_{\theta}(\mathbf{y}), \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y})) \right]$$

- $\mathcal{W}$  quantifies similarity between  $\mathcal{G}_{\theta}(y)$  and  $\pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$ .
- **Generator:**  $\mathcal{G}_{\theta}(y)$  is a probability distribution on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}}(y) \approx \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$ .

# Conditional Wasserstein GAN

**Input:** Supervised data  $(x_i, y_i) \in X \times Y$  generated by  $(\mathbf{x}, \mathbf{y}) \sim \mu$ .

**Task:** Sample from  $(\mathbf{x} \mid \mathbf{y} = y) \sim \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$  where posterior is unknown.

**Method:** Learn how to sample from posterior by solving

$$\hat{\theta} \in \arg \min_{\theta} \mathbb{E}_{\mathbf{y}} \left[ \mathcal{W}(\mathcal{G}_{\theta}(\mathbf{y}), \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y})) \right]$$

- $\mathcal{W}$  quantifies similarity between  $\mathcal{G}_{\theta}(y)$  and  $\pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$ .
- **Generator:**  $\mathcal{G}_{\theta}(y)$  is a probability distribution on  $X$ .
- **Desired property:**  $\mathcal{G}_{\hat{\theta}}(y) \approx \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$ .

Repeating the same procedure as for Wasserstein GAN ...

# Conditional Wasserstein GAN

**Input:** Supervised data  $(x_i, y_i) \in X \times Y$  generated by  $(\mathbf{x}, \mathbf{y}) \sim \mu$ .

**Task:** Sample from  $(\mathbf{x} \mid \mathbf{y} = y) \sim \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$  where posterior is unknown.

**Method:** Learn how to sample from posterior by solving

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \frac{1}{m} \sum_{i=1}^m \left[ D_{\phi}(x_i, y_i) - \mathbb{E}_{\mathbf{z} \sim \sigma} [D_{\phi}(G_{\theta}(\mathbf{z}, y_i), y_i)] \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \times Y \rightarrow \mathbb{R}$ .
- **Generator:**  $G_{\theta}: Z \times Y \rightarrow X$ .
- **Desired property:**  $G_{\theta}(\mathbf{z}, y)$  approximately  $\pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$ -distributed when  $\mathbf{z} \sim \sigma$ .

# Conditional Wasserstein GAN

**Input:** Supervised data  $(x_i, y_i) \in X \times Y$  generated by  $(\mathbf{x}, \mathbf{y}) \sim \mu$ .

**Task:** Sample from  $(\mathbf{x} \mid \mathbf{y} = y) \sim \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$  where posterior is unknown.

**Method:** Learn how to sample from posterior by solving

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \frac{1}{m} \sum_{i=1}^m \left[ D_{\phi}(x_i, y_i) - \mathbb{E}_{\mathbf{z} \sim \sigma} [D_{\phi}(G_{\theta}(\mathbf{z}, y_i), y_i)] \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \times Y \rightarrow \mathbb{R}$ .
- **Generator:**  $G_{\theta}: Z \times Y \rightarrow X$ .
- **Desired property:**  $G_{\theta}(\mathbf{z}, y)$  approximately  $\pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$ -distributed when  $\mathbf{z} \sim \sigma$ .
- **Issue:** Only single sample  $x_i \in X$  for each  $y_i \in Y$  in training data  
 $\implies G_{\hat{\theta}}(\mathbf{z}, y_i) \approx \delta_{x_i}$  independent of  $\mathbf{z} \sim \sigma$  (mode collapse).

# Conditional Wasserstein GAN

**Input:** Supervised data  $(x_i, y_i) \in X \times Y$  generated by  $(\mathbf{x}, \mathbf{y}) \sim \mu$ .

**Task:** Sample from  $(\mathbf{x} \mid \mathbf{y} = y) \sim \pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$  where posterior is unknown.

**Method:** Learn how to sample from posterior by solving

$$\hat{\theta} \in \arg \min_{\theta} \left\{ \max_{\phi} \left\{ \frac{1}{m} \sum_{i=1}^m \left[ D_{\phi}(x_i, y_i) - \mathbb{E}_{\mathbf{z} \sim \sigma} [D_{\phi}(G_{\theta}(\mathbf{z}, y_i), y_i)] \right] \right\} \right\}$$

- **Discriminator:**  $D_{\phi}: X \times Y \rightarrow \mathbb{R}$ .
- **Generator:**  $G_{\theta}: Z \times Y \rightarrow X$ .
- **Desired property:**  $G_{\theta}(\mathbf{z}, y)$  approximately  $\pi_{\text{post}}(\mathbf{x} \mid \mathbf{y} = y)$ -distributed when  $\mathbf{z} \sim \sigma$ .
- **Issue:** Only single sample  $x_i \in X$  for each  $y_i \in Y$  in training data  
 $\implies G_{\hat{\theta}}(\mathbf{z}, y_i) \approx \delta_{x_i}$  independent of  $\mathbf{z} \sim \sigma$  (mode collapse).

**Solution:** Allow discriminator to distinguish between unordered pairs of model parameter or random samples generated by generator.



- Adler, J., Lunz, S., Verdier, O., Schönlieb, C.-B., & Öktem, O. (2018a). Task adapted reconstruction for inverse problems. *ArXiv e-print*, *cs.CV*(1809.00948).
- Adler, J., Lunz, S., Verdier, O., Schönlieb, C.-B., & Öktem, O. (2018b). Task adapted reconstruction for inverse problems. In *NIPS 2018 meets medical imaging: Workshop within the 32nd annual conference on neural information processing systems (NIPS 2018)*. (Also available as ArXiv e-print: <https://arxiv.org/abs/1809.00948>)
- Adler, J., & Öktem, O. (2017). Solving ill-posed inverse problems using iterative deep neural networks. *Inverse Problems*, 33(12), 124007 (24pp). (Special issue: 100 Years of the Radon transform)
- Adler, J., & Öktem, O. (2018a). Deep Bayesian inversion: Computational uncertainty quantification for large scale inverse problems. *ArXiv e-print*, *stat.ML*(1811.05910).

## References II

- Adler, J., & Öktem, O. (2018b). Learned primal-dual reconstruction. *IEEE Transactions on Medical Imaging*, 37(6), 1322–1332.
- Arjovsky, M., Chintala, S., & Bottou, L. (2017). Wasserstein generative adversarial networks. *Proceedings of Machine Learning Research*, 70, 214–223. (Proceedings of the 34th International Conference on Machine Learning)
- Banert, S., Ringh, A., Adler, J., Karlsson, J., & Öktem, O. (2018). Data-driven nonsmooth optimization. *ArXiv*, 1808.00946.
- Dalca, A. V., Balakrishnan, G., Gutttag, J., & Sabuncu, M. R. (2018). Unsupervised learning for fast probabilistic diffeomorphic registration. In F. A. F., J. A. Schnabel, C. Davatzikos, L. C. Alberola, & G. Fichtinger (Eds.), *21st international conference on medical image computing and computer assisted intervention MICCAI 2018* (Vol. 11070, pp. 729–738). Springer Verlag.

## References III

- Farabet, C., Couprie, C., Najman, L., & LeCun, Y. (2013). Learning hierarchical features for scene labeling. *IEEE transactions on pattern analysis and machine intelligence*, 35(8), 1915–1929.
- Ghosal, S., & Ray, N. (2017). Deep deformable registration: Enhancing accuracy by fully convolutional neural net. *Pattern Recognition Letters*, 94, 81–86.
- Giryes, R., Eldar, Y. C., Bronstein, A. M., & Sapiro, G. (2017). Tradeoffs between convergence speed and reconstruction accuracy in inverse problems. *ArXiv e-print, cs.NA(1605.09232)*.
- Gregor, K., & LeCun, Y. (2010). Learning fast approximations of sparse coding. In *Proceedings of the 27th International Conference on International Conference on Machine Learning (ICML'10)* (pp. 399–406).
- Guo, Y., Liu, Y., Georgiou, T., & Lew, M. S. (2018). A review of semantic segmentation using deep neural networks. *International Journal of Multimedia Information Retrieval*, 7(2), 87–93.

## References IV

- Hammernik, K., Klatzer, E., T. ans Kobler, Recht, M. P., Sodickson, D. K., Pock, T., & Knoll, F. (2018). Learning a variational network for reconstruction of accelerated MRI data. *Magnetic Resonance in Medicine*, 79(6), 3055–3071.
- He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 770–778).
- Hsieh, J.-T., Zhao, S., Eismann, S., Mirabella, L., & Ermon, S. (2019). Learning neural PDE solvers with convergence guarantees. In *Seventh international conference on learning representations (ICLR 2019)*.
- Karpathy, A., & Fei-Fei, L. (2017). Deep visual-semantic alignments for generating image descriptions. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39(4), 664–676.

## References V

- Li, C. Y., Liang, X., Hu, Z., & Xing, E. P. (2018). Hybrid retrieval-generation reinforced agent for medical image report generation. *ArXiv e-print, cs.CV(1805.08298)*.
- Li, H., Schwab, J., Antholzer, S., & Haltmeier, M. (2018). NETT: Solving inverse problems with deep neural networks. *ArXiv, 1803.00092*.
- Lunz, S., Öktem, O., & Schönlieb, C.-B. (2018). Adversarial regularizers in inverse problems. *ArXiv, 1805.11572*. (NIPS 2018)
- Mao, X., Li, Q., Xie, H., Lau, R. Y., Wang, Z., & Smolley, S. P. (2016). Least squares generative adversarial networks. *ArXiv, 1611.04076*.
- Mardani, M., Gong, E., Cheng, J. Y., Vasanawala, S. S., Zaharchuk, G., Xing, L., & Pauly, J. M. (2019). Deep generative adversarial neural networks for compressive sensing MRI. *IEEE Transactions on Medical Imaging, 38*(1), 167–179.
- Oymak, S., & Soltanolkotabi, M. (2017). Fast and reliable parameter estimation from nonlinear observations. *SIAM Journal on Optimization, 27*(4), 2276–2300.

## References VI

- Putzky, P., & Welling, M. (2017). *Recurrent inference machines for solving inverse problems*. Retrieved from <https://openreview.net/pdf?id=HkS0lP9lg>
- Rizzuti, G., Siahkoohi, A., & Herrmann, F. J. (2019). *Learned iterative solvers for the Helmholtz equation*. Submitted to 81st EAGE Conference & Exhibition 2019. (Available from <https://www.slim.eos.ubc.ca/content/learned-iterative-solvers-helmholtz-equation>)
- Romano, Y., Elad, M., & Milanfar, P. (2017). The little engine that could: Regularization by denoising (RED). *SIAM Journal on Imaging Sciences*, 10(4), 1804–1844.
- Romano, Y., Isidoro, J., & Milanfar, P. (2017). RAISR: Rapid and accurate image super resolution. *IEEE Transactions on Computational Imaging*, 3(1), 110–125.
- Schwab, J., Antholzer, S., & Haltmeier, M. (2018). Deep null space learning for inverse problems: Convergence analysis and rates. *ArXiv*, 1806.06137.

## References VII

- Sermanet, P., Eigen, D., Zhang, X., Mathieu, M., Fergus, R., & LeCun, Y. (2013). OverFeat: Integrated recognition, localization and detection using convolutional networks. *ArXiv e-print, cs.CV(1312.6229)*.
- Syu, N.-S., Chen, Y.-S., & Chuang, Y.-Y. (2018). Learning deep convolutional networks for demosaicing. *ArXiv e-print, cs.CV(1802.03769)*.
- Thoma, M. (2016). A survey of semantic segmentation. *ArXiv e-print, cs.CV(1602.06541)*.
- Wolterink, J. M., Dinkla, A. M., Savenije, M. H. F., Seevinck, P. R., van den Berg, C. A. T., & Išgum, I. (2017). Deep MR to CT synthesis using unpaired data. In S. Tsafaris, A. Gooya, A. Frangi, & J. Prince (Eds.), *Simulation and synthesis in medical imaging. SASHIMI 2017* (Vol. 10557, pp. 14–23).
- Xie, J., Xu, L., & Chen, E. (2012). Image denoising and inpainting with deep neural networks. In *Proceedings of the 25th International Conference on Neural Information Processing Systems (NIPS 2012)* (pp. 341–349).

## References VIII

Zhu, B., Liu, J. Z., Cauley, S. F., Rosen, B. R., & Rosen, M. S. (2018). Image reconstruction by domain-transform manifold learning. *Nature*, 555, 487–492.